

# introduction to Data Analytics

July 17, 2021

## 1 Introduction Notebook

Estimated time needed: 10 minutes

### 1.1 Objectives

After completing this lab you will be able to:

- Acquire data in various ways
- Obtain insights from data with Pandas library

Table of Contents

Data Acquisition

Basic Insight of Dataset

Data Acquisition

There are various formats for a dataset: .csv, .json, .xlsx etc. The dataset can be stored in different places, on your local machine or sometimes online.

In this section, you will learn how to load a dataset into our Jupyter Notebook.

In our case, the Automobile Dataset is an online source, and it is in a CSV (comma separated value) format. Let's use this dataset as an example to practice data reading.

Data source: <https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>

Data type: csv

The Pandas Library is a useful tool that enables us to read various datasets into a dataframe; our Jupyter notebook platforms have a built-in Pandas Library so that all we need to do is import Pandas without installing.

```
[1]: # import pandas library
import pandas as pd
import numpy as np
```

Read Data

We use `pandas.read_csv()` function to read the csv file. In the brackets, we put the file path along with a quotation mark so that pandas will read the file into a dataframe from that address. The file path can be either an URL or your local file address.

Because the data does not include headers, we can add an argument `headers = None` inside the `read_csv()` method so that pandas will not automatically set the first row as a header.

You can also assign the dataset to any variable you create.

This dataset was hosted on IBM Cloud object. Click [HERE](#) for free storage.

```
[2]: # Import pandas library
import pandas as pd

# Read the online file by the URL provides above, and assign it to variable "df"
other_path = "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/
→IBMDeveloperSkillsNetwork-DA0101EN-SkillsNetwork/labs/Data%20files/auto.csv"
df = pd.read_csv(other_path, header=None)
```

After reading the dataset, we can use the `dataframe.head(n)` method to check the top n rows of the dataframe, where n is an integer. Contrary to `dataframe.head(n)`, `dataframe.tail(n)` will show you the bottom n rows of the dataframe.

```
[3]: # show the first 5 rows using dataframe.head() method
print("The first 5 rows of the dataframe")
df.head(5)
```

The first 5 rows of the dataframe

```
[3]:    0    1    2    3    4    5    6    7    8    9    ...  \
0    3    ?  alfa-romero  gas  std    two  convertible  rwd  front  88.6  ...
1    3    ?  alfa-romero  gas  std    two  convertible  rwd  front  88.6  ...
2    1    ?  alfa-romero  gas  std    two    hatchback  rwd  front  94.5  ...
3    2  164      audi  gas  std    four      sedan  fwd  front  99.8  ...
4    2  164      audi  gas  std    four      sedan  4wd  front  99.4  ...
```

```
      16    17    18    19    20    21    22  23  24    25
0  130  mpfi  3.47  2.68   9.0  111  5000  21  27  13495
1  130  mpfi  3.47  2.68   9.0  111  5000  21  27  16500
2  152  mpfi  2.68  3.47   9.0  154  5000  19  26  16500
3  109  mpfi  3.19  3.40  10.0  102  5500  24  30  13950
4  136  mpfi  3.19  3.40   8.0  115  5500  18  22  17450
```

```
[5 rows x 26 columns]
```

Question #1:

Check the bottom 10 rows of data frame "df".

```
[4]: # Write your code below and press Shift+Enter to execute
df.tail(10)
```

```
[4]:      0      1      2      3      4      5      6      7      8      9      ...      16  \
195  -1     74  volvo    gas    std   four  wagon  rwd  front  104.3  ...   141
196  -2    103  volvo    gas    std   four  sedan  rwd  front  104.3  ...   141
197  -1     74  volvo    gas    std   four  wagon  rwd  front  104.3  ...   141
198  -2    103  volvo    gas  turbo   four  sedan  rwd  front  104.3  ...   130
199  -1     74  volvo    gas  turbo   four  wagon  rwd  front  104.3  ...   130
200  -1     95  volvo    gas    std   four  sedan  rwd  front  109.1  ...   141
201  -1     95  volvo    gas  turbo   four  sedan  rwd  front  109.1  ...   141
202  -1     95  volvo    gas    std   four  sedan  rwd  front  109.1  ...   173
203  -1     95  volvo  diesel  turbo   four  sedan  rwd  front  109.1  ...   145
204  -1     95  volvo    gas  turbo   four  sedan  rwd  front  109.1  ...   141
```

```
      17      18      19      20      21      22      23      24      25
195  mpfi  3.78  3.15   9.5  114  5400  23  28  13415
196  mpfi  3.78  3.15   9.5  114  5400  24  28  15985
197  mpfi  3.78  3.15   9.5  114  5400  24  28  16515
198  mpfi  3.62  3.15   7.5  162  5100  17  22  18420
199  mpfi  3.62  3.15   7.5  162  5100  17  22  18950
200  mpfi  3.78  3.15   9.5  114  5400  23  28  16845
201  mpfi  3.78  3.15   8.7  160  5300  19  25  19045
202  mpfi  3.58  2.87   8.8  134  5500  18  23  21485
203  idi   3.01  3.40  23.0  106  4800  26  27  22470
204  mpfi  3.78  3.15   9.5  114  5400  19  25  22625
```

[10 rows x 26 columns]

Click here for the solution

```
print("The last 10 rows of the dataframe\n")
df.tail(10)
```

Add Headers

Take a look at our dataset. Pandas automatically set the header with an integer starting from 0.

To better describe our data, we can introduce a header. This information is available at: <https://archive.ics.uci.edu/ml/datasets/Automobile>.

Thus, we have to add headers manually.

First, we create a list “headers” that include all column names in order. Then, we use `dataframe.columns = headers` to replace the headers with the list we created.

```
[5]: # create headers list
headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "
→ "num-of-doors", "body-style",
```

```

        "drive-wheels","engine-location","wheel-base",\
        →"length","width","height","curb-weight","engine-type",
        "num-of-cylinders",\
        →"engine-size","fuel-system","bore","stroke","compression-ratio","horsepower",
        "peak-rpm","city-mpg","highway-mpg","price"]
print("headers\n", headers)

```

headers

```

['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration', 'num-of-
doors', 'body-style', 'drive-wheels', 'engine-location', 'wheel-base', 'length',
'width', 'height', 'curb-weight', 'engine-type', 'num-of-cylinders', 'engine-
size', 'fuel-system', 'bore', 'stroke', 'compression-ratio', 'horsepower',
'peak-rpm', 'city-mpg', 'highway-mpg', 'price']

```

We replace headers and recheck our dataframe:

```

[6]: df.columns = headers
df.head(10)

```

```

[6]:
symboling normalized-losses      make fuel-type aspiration num-of-doors \
0          3          ?  alfa-romero    gas      std         two
1          3          ?  alfa-romero    gas      std         two
2          1          ?  alfa-romero    gas      std         two
3          2        164    audi        gas      std         four
4          2        164    audi        gas      std         four
5          2          ?    audi        gas      std         two
6          1        158    audi        gas      std         four
7          1          ?    audi        gas      std         four
8          1        158    audi        gas    turbo         four
9          0          ?    audi        gas    turbo         two

      body-style drive-wheels engine-location wheel-base ... engine-size \
0  convertible      rwd      front      88.6 ...      130
1  convertible      rwd      front      88.6 ...      130
2   hatchback      rwd      front      94.5 ...      152
3      sedan      fwd      front      99.8 ...      109
4      sedan      4wd      front      99.4 ...      136
5      sedan      fwd      front      99.8 ...      136
6      sedan      fwd      front     105.8 ...      136
7      wagon      fwd      front     105.8 ...      136
8      sedan      fwd      front     105.8 ...      131
9   hatchback      4wd      front      99.5 ...      131

      fuel-system bore  stroke compression-ratio horsepower  peak-rpm city-mpg \
0          mpfi  3.47   2.68             9.0          111    5000    21
1          mpfi  3.47   2.68             9.0          111    5000    21
2          mpfi  2.68   3.47             9.0          154    5000    19
3          mpfi  3.19   3.40            10.0          102    5500    24

```

4	mpfi	3.19	3.40	8.0	115	5500	18
5	mpfi	3.19	3.40	8.5	110	5500	19
6	mpfi	3.19	3.40	8.5	110	5500	19
7	mpfi	3.19	3.40	8.5	110	5500	19
8	mpfi	3.13	3.40	8.3	140	5500	17
9	mpfi	3.13	3.40	7.0	160	5500	16

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450
5	25	15250
6	25	17710
7	25	18920
8	20	23875
9	22	?

[10 rows x 26 columns]

We need to replace the “?” symbol with NaN so the dropna() can remove the missing values:

```
[7]: df1=df.replace('?',np.NaN)
```

We can drop missing values along the column “price” as follows:

```
[8]: df=df1.dropna(subset=["price"], axis=0)
df.head(20)
```

```
[8]:      symboling normalized-losses      make fuel-type aspiration \
0          3          NaN  alfa-romero    gas    std
1          3          NaN  alfa-romero    gas    std
2          1          NaN  alfa-romero    gas    std
3          2         164      audi    gas    std
4          2         164      audi    gas    std
5          2          NaN      audi    gas    std
6          1         158      audi    gas    std
7          1          NaN      audi    gas    std
8          1         158      audi    gas  turbo
10         2         192      bmw    gas    std
11         0         192      bmw    gas    std
12         0         188      bmw    gas    std
13         0         188      bmw    gas    std
14         1          NaN      bmw    gas    std
15         0          NaN      bmw    gas    std
16         0          NaN      bmw    gas    std
```

17	0	NaN	bmw	gas	std
18	2	121	chevrolet	gas	std
19	1	98	chevrolet	gas	std
20	0	81	chevrolet	gas	std

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
0	two	convertible	rwd	front	88.6	...	
1	two	convertible	rwd	front	88.6	...	
2	two	hatchback	rwd	front	94.5	...	
3	four	sedan	fwd	front	99.8	...	
4	four	sedan	4wd	front	99.4	...	
5	two	sedan	fwd	front	99.8	...	
6	four	sedan	fwd	front	105.8	...	
7	four	wagon	fwd	front	105.8	...	
8	four	sedan	fwd	front	105.8	...	
10	two	sedan	rwd	front	101.2	...	
11	four	sedan	rwd	front	101.2	...	
12	two	sedan	rwd	front	101.2	...	
13	four	sedan	rwd	front	101.2	...	
14	four	sedan	rwd	front	103.5	...	
15	four	sedan	rwd	front	103.5	...	
16	two	sedan	rwd	front	103.5	...	
17	four	sedan	rwd	front	110.0	...	
18	two	hatchback	fwd	front	88.4	...	
19	two	hatchback	fwd	front	94.5	...	
20	four	sedan	fwd	front	94.5	...	

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
0	130	mpfi	3.47	2.68	9.0	111	
1	130	mpfi	3.47	2.68	9.0	111	
2	152	mpfi	2.68	3.47	9.0	154	
3	109	mpfi	3.19	3.40	10.0	102	
4	136	mpfi	3.19	3.40	8.0	115	
5	136	mpfi	3.19	3.40	8.5	110	
6	136	mpfi	3.19	3.40	8.5	110	
7	136	mpfi	3.19	3.40	8.5	110	
8	131	mpfi	3.13	3.40	8.3	140	
10	108	mpfi	3.50	2.80	8.8	101	
11	108	mpfi	3.50	2.80	8.8	101	
12	164	mpfi	3.31	3.19	9.0	121	
13	164	mpfi	3.31	3.19	9.0	121	
14	164	mpfi	3.31	3.19	9.0	121	
15	209	mpfi	3.62	3.39	8.0	182	
16	209	mpfi	3.62	3.39	8.0	182	
17	209	mpfi	3.62	3.39	8.0	182	
18	61	2bbl	2.91	3.03	9.5	48	
19	90	2bbl	3.03	3.11	9.6	70	

20                    90                    2bb1 3.03        3.11                    9.6                    70

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500
3	5500	24	30	13950
4	5500	18	22	17450
5	5500	19	25	15250
6	5500	19	25	17710
7	5500	19	25	18920
8	5500	17	20	23875
10	5800	23	29	16430
11	5800	23	29	16925
12	4250	21	28	20970
13	4250	21	28	21105
14	4250	20	25	24565
15	5400	16	22	30760
16	5400	16	22	41315
17	5400	15	20	36880
18	5100	47	53	5151
19	5400	38	43	6295
20	5400	38	43	6575

[20 rows x 26 columns]

Now, we have successfully read the raw dataset and added the correct headers into the dataframe.

Question #2:

Find the name of the columns of the dataframe.

```
[9]: # Write your code below and press Shift+Enter to execute
print(df.columns)
```

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
      'highway-mpg', 'price'],
      dtype='object')
```

[Click here for the solution](#)

```
print(df.columns)
```

Save Dataset

Correspondingly, Pandas enables us to save the dataset to csv. By using the `dataframe.to_csv()` method, you can add the file path and name along with quotation marks in the brackets.

For example, if you would save the dataframe `df` as `automobile.csv` to your local machine, you may use the syntax below, where `index = False` means the row names will not be written. `df.to_csv("automobile.csv", index = False)` We can also read and save other file formats. We can use similar functions like `pd.read_csv()` and `df.to_csv()` for other formats.

## Read/Save Other Data Formats

Data Format	Read	Save
csv	<code>pd.read_csv()</code>	<code>df.to_csv()</code>
json	<code>pd.read_json()</code>	<code>df.to_json()</code>
excel	<code>pd.read_excel()</code>	<code>df.to_excel()</code>
hdf	<code>pd.read_hdf()</code>	<code>df.to_hdf()</code>
sql	<code>pd.read_sql()</code>	<code>df.to_sql()</code>
...	...	...

## Basic Insight of Dataset

After reading data into Pandas dataframe, it is time for us to explore the dataset.

There are several ways to obtain essential insights of the data to help us better understand our dataset.

## Data Types

Data has a variety of types.

The main types stored in Pandas dataframes are object, float, int, bool and datetime64. In order to better learn about each attribute, it is always good for us to know the data type of each column. In Pandas:

```
[10]: df.dtypes
```

```
[10]: symboling          int64
normalized-losses      object
make                   object
fuel-type              object
aspiration             object
num-of-doors           object
body-style             object
drive-wheels           object
engine-location        object
wheel-base            float64
length                float64
width                  float64
height                 float64
curb-weight            int64
engine-type            object
num-of-cylinders       object
engine-size            int64
fuel-system            object
```



```
bore          object
stroke        object
compression-ratio float64
horsepower    object
peak-rpm      object
city-mpg      int64
highway-mpg   int64
price         object
dtype: object
```

A series with the data type of each column is returned.

```
[11]: # check the data type of data frame "df" by .dtypes
print(df.dtypes)
```

```
symboling      int64
normalized-losses object
make           object
fuel-type      object
aspiration     object
num-of-doors   object
body-style     object
drive-wheels   object
engine-location object
wheel-base    float64
length         float64
width          float64
height         float64
curb-weight    int64
engine-type    object
num-of-cylinders object
engine-size    int64
fuel-system    object
bore           object
stroke         object
compression-ratio float64
horsepower     object
peak-rpm       object
city-mpg       int64
highway-mpg    int64
price          object
dtype: object
```

As shown above, it is clear to see that the data type of “symboling” and “curb-weight” are int64, “normalized-losses” is object, and “wheel-base” is float64, etc.

These data types can be changed; we will learn how to accomplish this in a later module.

Describe

If we would like to get a statistical summary of each column e.g. count, column mean value, column standard deviation, etc., we use the describe method: `dataframe.describe()` This method will provide various summary statistics, excluding NaN (Not a Number) values.

```
[12]: df.describe()
```

```
[12]:
```

	symboling	wheel-base	length	width	height	\
count	201.000000	201.000000	201.000000	201.000000	201.000000	
mean	0.840796	98.797015	174.200995	65.889055	53.766667	
std	1.254802	6.066366	12.322175	2.101471	2.447822	
min	-2.000000	86.600000	141.100000	60.300000	47.800000	
25%	0.000000	94.500000	166.800000	64.100000	52.000000	
50%	1.000000	97.000000	173.200000	65.500000	54.100000	
75%	2.000000	102.400000	183.500000	66.600000	55.500000	
max	3.000000	120.900000	208.100000	72.000000	59.800000	

  

	curb-weight	engine-size	compression-ratio	city-mpg	highway-mpg	
count	201.000000	201.000000	201.000000	201.000000	201.000000	
mean	2555.666667	126.875622	10.164279	25.179104	30.686567	
std	517.296727	41.546834	4.004965	6.423220	6.815150	
min	1488.000000	61.000000	7.000000	13.000000	16.000000	
25%	2169.000000	98.000000	8.600000	19.000000	25.000000	
50%	2414.000000	120.000000	9.000000	24.000000	30.000000	
75%	2926.000000	141.000000	9.400000	30.000000	34.000000	
max	4066.000000	326.000000	23.000000	49.000000	54.000000	

This shows the statistical summary of all numeric-typed (int, float) columns.

For example, the attribute “symboling” has 205 counts, the mean value of this column is 0.83, the standard deviation is 1.25, the minimum value is -2, 25th percentile is 0, 50th percentile is 1, 75th percentile is 2, and the maximum value is 3.

However, what if we would also like to check all the columns including those that are of type object?

You can add an argument `include = “all”` inside the bracket. Let’s try it again.

```
[13]: # describe all the columns in "df"
df.describe(include = "all")
```

```
[13]:
```

	symboling	normalized-losses	make	fuel-type	aspiration	\
count	201.000000	164	201	201	201	
unique	NaN	51	22	2	2	
top	NaN	161	toyota	gas	std	
freq	NaN	11	32	181	165	
mean	0.840796	NaN	NaN	NaN	NaN	
std	1.254802	NaN	NaN	NaN	NaN	
min	-2.000000	NaN	NaN	NaN	NaN	
25%	0.000000	NaN	NaN	NaN	NaN	

50%	1.000000	NaN	NaN	NaN	NaN
75%	2.000000	NaN	NaN	NaN	NaN
max	3.000000	NaN	NaN	NaN	NaN

  

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	\
count	199	201	201	201	201.000000	...	
unique	2	5	3	2	NaN	...	
top	four	sedan	fwd	front	NaN	...	
freq	113	94	118	198	NaN	...	
mean	NaN	NaN	NaN	NaN	98.797015	...	
std	NaN	NaN	NaN	NaN	6.066366	...	
min	NaN	NaN	NaN	NaN	86.600000	...	
25%	NaN	NaN	NaN	NaN	94.500000	...	
50%	NaN	NaN	NaN	NaN	97.000000	...	
75%	NaN	NaN	NaN	NaN	102.400000	...	
max	NaN	NaN	NaN	NaN	120.900000	...	

  

	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower	\
count	201.000000	201	197	197	201.000000	199	
unique	NaN	8	38	36	NaN	58	
top	NaN	mpfi	3.62	3.40	NaN	68	
freq	NaN	92	23	19	NaN	19	
mean	126.875622	NaN	NaN	NaN	10.164279	NaN	
std	41.546834	NaN	NaN	NaN	4.004965	NaN	
min	61.000000	NaN	NaN	NaN	7.000000	NaN	
25%	98.000000	NaN	NaN	NaN	8.600000	NaN	
50%	120.000000	NaN	NaN	NaN	9.000000	NaN	
75%	141.000000	NaN	NaN	NaN	9.400000	NaN	
max	326.000000	NaN	NaN	NaN	23.000000	NaN	

  

	peak-rpm	city-mpg	highway-mpg	price
count	199	201.000000	201.000000	201
unique	22	NaN	NaN	186
top	5500	NaN	NaN	16500
freq	36	NaN	NaN	2
mean	NaN	25.179104	30.686567	NaN
std	NaN	6.423220	6.815150	NaN
min	NaN	13.000000	16.000000	NaN
25%	NaN	19.000000	25.000000	NaN
50%	NaN	24.000000	30.000000	NaN
75%	NaN	30.000000	34.000000	NaN
max	NaN	49.000000	54.000000	NaN

[11 rows x 26 columns]

Now it provides the statistical summary of all the columns, including object-typed attributes.

We can now see how many unique values there, which one is the top value and the frequency of

top value in the object-typed columns.

Some values in the table above show as “NaN”. This is because those numbers are not available regarding a particular column type.

Question #3:

You can select the columns of a dataframe by indicating the name of each column. For example, you can select the three columns as follows:

```
dataframe[['column 1','column 2','column 3']]
```

Where “column” is the name of the column, you can apply the method “.describe()” to get the statistics of those columns as follows:

```
dataframe[['column 1','column 2','column 3']].describe()
```

Apply the method to “.describe()” to the columns ‘length’ and ‘compression-ratio’.

```
[14]: # Write your code below and press Shift+Enter to execute
df[['length','compression-ratio']].describe()
```

```
[14]:
```

	length	compression-ratio
count	201.000000	201.000000
mean	174.200995	10.164279
std	12.322175	4.004965
min	141.100000	7.000000
25%	166.800000	8.600000
50%	173.200000	9.000000
75%	183.500000	9.400000
max	208.100000	23.000000

[Click here for the solution](#)

```
df[['length','compression-ratio']].describe()
```

Info

Another method you can use to check your dataset is: dataframe.info() It provides a concise summary of your DataFrame.

This method prints information about a DataFrame including the index dtype and columns, non-null values and memory usage.

```
[15]: # look at the info of "df"
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 201 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              201 non-null    int64
1   normalized-losses      164 non-null    object
```

2	make	201 non-null	object
3	fuel-type	201 non-null	object
4	aspiration	201 non-null	object
5	num-of-doors	199 non-null	object
6	body-style	201 non-null	object
7	drive-wheels	201 non-null	object
8	engine-location	201 non-null	object
9	wheel-base	201 non-null	float64
10	length	201 non-null	float64
11	width	201 non-null	float64
12	height	201 non-null	float64
13	curb-weight	201 non-null	int64
14	engine-type	201 non-null	object
15	num-of-cylinders	201 non-null	object
16	engine-size	201 non-null	int64
17	fuel-system	201 non-null	object
18	bore	197 non-null	object
19	stroke	197 non-null	object
20	compression-ratio	201 non-null	float64
21	horsepower	199 non-null	object
22	peak-rpm	199 non-null	object
23	city-mpg	201 non-null	int64
24	highway-mpg	201 non-null	int64
25	price	201 non-null	object

dtypes: float64(5), int64(5), object(16)

memory usage: 42.4+ KB

Excellent! You have just completed the Introduction Notebook!

### 1.1.1 Thank you for completing this lab!

## 1.2 Author

Joseph Santarcangelo

### 1.2.1 Other Contributors

Mahdi Noorian PhD

Bahare Talayian

Eric Xiao

Steven Dong

Parizad

Hima Vasudevan

Fiorella Wenver

Yi Yao.

### 1.3 Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-10-30	2.3	Lakshmi	Changed URL of the csv
2020-09-22	2.2	Nayef	Added replace() method to remove '?'
2020-09-09	2.1	Lakshmi	Made changes in info method of dataframe
2020-08-27	2.0	Lavanya	Moved lab to course repo in GitLab

##

© IBM Corporation 2020. All rights reserved.